



EXTREME GRADIENT BOOSTING WITH XGBOOST

Review of pipelines using sklearn

Pipeline Review

- Takes a list of named 2-tuples (name, pipeline_step) as input
- Tuples can contain any arbitrary scikit-learn compatible estimator or transformer object
- Pipeline implements fit/predict methods
- Can be used as input estimator into grid/randomized search and cross_val_score methods

Scikit-learn pipeline example

```
In [1]: import pandas as pd
...: from sklearn.ensemble import RandomForestRegressor
...: import numpy as np
...: from sklearn.preprocessing import StandardScaler
...: from sklearn.pipeline import Pipeline
...: from sklearn.model_selection import cross_val_score

In [2]: names = ["crime", "zone", "industry", "charles",
...: "no", "rooms", "age", "distance",
...: "radial", "tax", "pupil", "aam", "lower", "med_price"]

In [3]: data = pd.read_csv("boston_housing.csv", names=names)

In [4]: X, y = data.iloc[:, :-1], data.iloc[:, -1]

In [5]: rf_pipeline = Pipeline(["st_scaler",
...: StandardScaler(),
...: ("rf_model", RandomForestRegressor())]

In [6]: scores = cross_val_score(rf_pipeline, X, y,
```

Scikit-learn pipeline example

```
In [7]: final_avg_rmse = np.mean(np.sqrt(np.abs(scores)))
```

```
In [8]: print("Final RMSE:", final_avg_rmse)  
Final RMSE: 4.54530686529
```

Preprocessing I: LabelEncoder and OneHotEncoder

- LabelEncoder: Converts a categorical column of strings into integers
- OneHotEncoder: Takes the column of integers and encodes them as dummy variables
- Cannot be done within a pipeline

Preprocessing II: DictVectorizer

- Traditionally used in text processing
- Converts lists of feature mappings into vectors
- Need to convert DataFrame into a list of dictionary entries
- Explore the [scikit-learn documentation](#)



EXTREME GRADIENT BOOSTING WITH XGBOOST

Let's build pipelines!



EXTREME GRADIENT BOOSTING WITH XGBOOST

Incorporating xgboost into pipelines

Scikit-Learn Pipeline Example With XGBoost

```
In [1]: import pandas as pd
...: import xgboost as xgb
...: import numpy as np
...: from sklearn.preprocessing import StandardScaler
...: from sklearn.pipeline import Pipeline
...: from sklearn.model_selection import cross_val_score

In [2]: names = ["crime", "zone", "industry", "charles", "no",
...: "rooms", "age", "distance", "radial", "tax",
...: "pupil", "aam", "lower", "med_price"]

In [3]: data = pd.read_csv("boston_housing.csv", names=names)
In [4]: X, y = data.iloc[:, :-1], data.iloc[:, -1]

In [5]: xgb_pipeline = Pipeline(["st_scaler",
...: StandardScaler()),
...: ("xgb_model", xgb.XGBRegressor())]

In [6]: scores = cross_val_score(xgb_pipeline, X, y,
...: scoring="neg_mean_squared_error", cv=10)

In [7]: final_avg_rmse = np.mean(np.sqrt(np.abs(scores)))
In [8]: print("Final XGB RMSE: ", final_avg_rmse)
```

Additional Components Introduced For Pipelines

- `sklearn_pandas`:
 - `DataFrameMapper` - Interoperability between pandas and scikit-learn
 - `CategoricalImputer` - Allow for imputation of categorical variables before conversion to integers
- `sklearn.preprocessing`:
 - `Imputer` - Native imputation of numerical columns in scikit-learn
- `sklearn.pipeline`:
 - `FeatureUnion` - combine multiple pipelines of features into a single



EXTREME GRADIENT BOOSTING WITH XGBOOST

Let's practice!



EXTREME GRADIENT BOOSTING WITH XGBOOST

Tuning xgboost hyperparameters in a pipeline

Sergey Fogelson
VP of Analytics, Vincom

Tuning XGBoost hyperparameters in a Pipeline

```
In [1]: import pandas as pd
...: import xgboost as xgb
...: import numpy as np
...: from sklearn.preprocessing import StandardScaler
...: from sklearn.pipeline import Pipeline
...: from sklearn.model_selection import RandomizedSearchCV

In [2]: names = ["crime", "zone", "industry", "charles", "no",
...:              "rooms", "age", "distance", "radial", "tax",
...:              "pupil", "aam", "lower", "med_price"]
In [3]: data = pd.read_csv("boston_housing.csv", names=names)
In [4]: X, y = data.iloc[:, :-1], data.iloc[:, -1]
In [5]: xgb_pipeline = Pipeline([("st_scaler",
...:                               StandardScaler()), ("xgb_model", xgb.XGBRegressor())])

In [6]: gbm_param_grid = {
...:     'xgb_model__subsample': np.arange(.05, 1, .05),
...:     'xgb_model__max_depth': np.arange(3, 20, 1),
...:     'xgb_model__colsample_bytree': np.arange(.1, 1.05, .05) }
```

```
In [7]: randomized_search_cv = RandomizedSearchCV(estimator=xgb_pipeline,
```

Tuning XGBoost hyperparameters in a Pipeline II

```
In [9]: print("Best rmse: ",
...: np.sqrt(np.abs(randomized_neg_mse.best_score_)))
Best rmse: 3.9966784203040677

In [10]: print("Best model: ",
...: randomized_neg_mse.best_estimator_)
Best model: Pipeline(steps=[('st_scaler', StandardScaler(copy=True,
with_mean=True, with_std=True)),
('xgb_model', XGBRegressor(base_score=0.5, colsample_bylevel=1,
colsample_bytree=0.95000000000000029, gamma=0, learning_rate=0.1,
max_delta_step=0, max_depth=8, min_child_weight=1, missing=None,
n_estimators=100, nthread=-1, objective='reg:linear', reg_alpha=0,
reg_lambda=1, scale_pos_weight=1, seed=0, silent=True,
subsample=0.90000000000000013))])
```



EXTREME GRADIENT BOOSTING WITH XGBOOST

Let's finish this up!



EXTREME GRADIENT BOOSTING WITH XGBOOST

Final Thoughts

What We Have Covered And You Have Learned

- Using XGBoost for classification tasks
- Using XGBoost for regression tasks
- Tuning XGBoost's most important hyperparameters
- Incorporating XGBoost into sklearn pipelines

What We Have Not Covered (And How You Can Proceed)

- Using XGBoost for ranking/recommendation problems (Netflix/Amazon problem)
- Using more sophisticated hyperparameter tuning strategies for tuning XGBoost models (Bayesian Optimization)
- Using XGBoost as part of an ensemble of other models for regression/classification



EXTREME GRADIENT BOOSTING WITH XGBOOST

Congratulations!